

The Metatheory of Crush’s Certified Lowering Pipeline

lean-crush

August 2026

Abstract

Crush translates a supported finite collection of Lean propositions into SMT commands. This document explains the current mathematical framework for that translation: intrinsically typed higher-order syntax, datatype-aware source models, total flattened defunctionalization, raw SMT command semantics, native datatype declarations, guarded carrier encodings, and exact agreement with a completed production command snapshot. It is written in dependency order: each object and notation is introduced before the first theorem that uses it. The final theorem justifies the reified higher-order theory and the exact semantic command set produced by the modeled lowering pipeline; it does not assume that an external solver is correct and does not assign a denotation to arbitrary `Lean.Expr`.

Contents

1	Reader’s guide	2
1.1	What is being proved	2
1.2	Proof stages	3
1.3	Metavariables and notation	3
2	The intrinsically typed higher-order source	4
2.1	Types, contexts, and signatures	4
2.2	Source models	5
3	Datatype-aware source semantics	5
3.1	Intrinsic datatype blocks	6
3.2	Finite free values	6
3.3	Lawful source models	6
4	The intrinsically typed first-order target	7
5	Total flattened defunctionalization	10
5.1	Flattened arrow shape	10
5.2	Captures, closures, and symbols	10
5.3	What translation returns	11
5.4	How terms are translated	11
6	The canonical model and flattened correctness	12
6.1	Modeling function values	12
6.2	The unary reference proof	14


7	Raw SMT semantics	14
8	Exact FO-to-SMT representation	15
8.1	Native datatype representation	16
8.2	One induced raw model	17
9	Guarded carriers and guarded SMT encoding	18
9.1	Guarded carrier equivalence	18
9.2	Datatype guard lifting	19
9.3	Guard-aware raw syntax	19
9.4	Native blocks and recursive guards in one model	20
10	Lean reification and intrinsic VCG	20
10.1	The supported Lean boundary	21
10.2	Pure aggregate VCG	21
11	Production certificates and whole-theory agreement	22
11.1	Exact final-state component traces	22
11.2	Whole-theory comparison	23
11.3	Operational trust status	23
12	Proof architecture and exact scope	23
12.1	What the composed theorem proves	24
12.2	Exact boundaries	25
13	Machine-checked theorem inventory	25

1 Reader's guide

Contents of this section. The question answered by the formalization; proof stages; notation and scope.

1.1 What is being proved

The formalization studies this pipeline:



The first arrow is a partial recognizer: every accepted fact is reified under one common signature and datatype environment, and unsupported Lean syntax is rejected. The remaining intrinsic transformations are total. Native datatype commands and recursive well-formedness definitions are modeled components of the final command set. The principal semantic conclusion is:

If the completed represented command set is semantically unsatisfiable, then the reified higher-order source theory is unsatisfiable in every source model satisfying the retained datatype laws.

The proof is by countermodel construction. Starting from one lawful model of the complete higher-order theory, Lean constructs one model of all flattened sentences and generated obligations, then constructs one relational SMT model satisfying native declarations, derived recursive guards, ordinary declarations, and assertions. Contraposition yields the desired reflection of unsatisfiability.

1.2 Proof stages

The argument is organized into seven stages:

1. Define an intrinsically typed higher-order source language and its ordinary function-space semantics.
2. Describe productive monomorphic datatype blocks, their finite free-algebra semantics, and the source-model lawfulness contract.
3. Define an intrinsically typed first-order target whose function-value sorts are opaque, then give total flattened defunctionalization.
4. Interpret every translated sentence in one canonical model and prove simultaneous validity of the aggregate target theory.
5. Give raw semantics to the admitted SMT command fragment and prove one exact command-set representation theorem.
6. Lift guarded and datatype carriers, validate native datatype and recursive guard commands, and combine all components in one raw model.
7. Retain structural Lean reification, exact stateful VCG execution, and proof-producing agreement with a completed production snapshot.

1.3 Metavariables and notation

The following metavariables are used consistently throughout the document. Semantic and translation notation is introduced later, beside the definition that gives it meaning.

e, f, a	source or target terms; f is usually function-valued,
φ, ψ	Boolean formulas; a closed formula is a sentence,
Φ	a finite higher-order source theory,
β	an opaque source base-sort name,
$\alpha, \gamma, \sigma, \tau, \nu$	source types,
κ	a first-order sort,
Γ, Δ	typing contexts,
Σ	a typed signature,
T	a typed first-order theory,
M, N	semantic models,
ρ	a source valuation or typed renaming,
ν	a target valuation,
v, w	individual semantic values,
$\mathcal{E}$	a concrete FO-to-SMT encoding choice,
$\mathcal{G}$	a guard-aware SMT encoding choice,
C	an ordered array of raw SMT commands.

Subscripts s , t , and q mean source, typed target, and raw SMT model, respectively. Proof-facing Lean declarations use the corresponding Greek binders; executable allocation and state code keeps descriptive names when several objects have the same mathematical role.

2 The intrinsically typed higher-order source

Contents of this section. Types, contexts, terms, and ordinary higher-order semantics.

2.1 Types, contexts, and signatures

Definition 2.1 (Source types). *Let β range over opaque base-sort names. Source types are generated by*

$$\tau ::= \text{Bool} \mid \text{Base}(\beta) \mid \tau \rightarrow \sigma.$$

*Arrows are nondependent and associate to the right. A context $\Gamma = [\tau_0, \dots, \tau_{n-1}]$ records local binders with the newest binder at the head. We write Γ, τ for extension by a new binder of type τ ; this is stored as $\tau :: \Gamma$ in Lean. A signature Σ is a list of the types of global constants.*¹

We write $\Gamma \ni \tau$ for the intrinsically typed lookup judgment. Its inhabitants are de Bruijn references, generated by:

$$\frac{}{Z : \Gamma, \tau \ni \tau} \quad \frac{i : \Gamma \ni \tau}{Si : \Gamma, \sigma \ni \tau}.$$

The judgment is intrinsically witnessed: its derivation simultaneously records the position and proves the type at that position. The witness Z denotes index zero and Si denotes the successor of the index denoted by i . Forgetting the derivation yields the natural-number index used when collecting free variables. Constants use the same typed-reference representation in Σ . Therefore a variable or constant with the wrong type cannot be constructed.

For $i : \Gamma \ni \tau$, write x_i for the local-variable term selected by i . For $k : \Sigma \ni \tau$, write c_k for the global-constant term selected by k from the signature. The subscripts are typed lookup witnesses, not variable or constant names.

Definition 2.2 (Source terms). *The intrinsic judgment $\Sigma; \Gamma \vdash e : \tau$ has variables, constants, Boolean literals and connectives, same-type equality, lambda, application, and universal and existential quantification:*

$$\frac{i : \Gamma \ni \tau}{\Sigma; \Gamma \vdash x_i : \tau} \quad \frac{k : \Sigma \ni \tau}{\Sigma; \Gamma \vdash c_k : \tau} \quad \frac{\Sigma; \Gamma, \tau \vdash e : \sigma}{\Sigma; \Gamma \vdash \lambda x:\tau. e : \tau \rightarrow \sigma} \quad \frac{\Sigma; \Gamma \vdash f : \tau \rightarrow \sigma \quad \Sigma; \Gamma \vdash a : \tau}{\Sigma; \Gamma \vdash f a : \sigma}$$

$$\frac{\Sigma; \Gamma \vdash e_1 : \tau \quad \Sigma; \Gamma \vdash e_2 : \tau}{\Sigma; \Gamma \vdash e_1 = e_2 : \text{Bool}} \quad \frac{\Sigma; \Gamma, \tau \vdash \varphi : \text{Bool} \quad Q \in \{\forall, \exists\}}{\Sigma; \Gamma \vdash Qx:\tau. \varphi : \text{Bool}}.$$

*The omitted rules are the standard typed Boolean rules. A formula is a term of type `Bool`; a sentence is a formula in the empty context.*²

Quantifiers may range over arrow types. Thus the modeled language is genuinely higher-order even though its arrows are nondependent.

¹Lean formalization: `Crush/Metattheory/H0/Core.lean`.

²Lean formalization: `Crush/Metattheory/H0/Core.lean`.

2.2 Source models

Definition 2.3 (Source type denotation). A carrier family B assigns a nonempty type $B(\beta)$ to each opaque base sort. Its extension to source types is

$$\llbracket \text{Bool} \rrbracket_B = \text{Prop}, \quad \llbracket \text{Base}(\beta) \rrbracket_B = B(\beta), \quad \llbracket \tau \rightarrow \sigma \rrbracket_B = \llbracket \tau \rrbracket_B \rightarrow \llbracket \sigma \rrbracket_B.$$

Thus the only choices in B are the carriers assigned to opaque base sorts.

Definition 2.4 (Source model and valuation). A source model M_s consists of a carrier family B and an interpretation $c_k^{M_s} \in \llbracket \tau \rrbracket_B$ for each $k : \Sigma \ni \tau$. A valuation ρ for Γ assigns a value in $\llbracket \tau \rrbracket_B$ to every $i : \Gamma \ni \tau$. We write $\rho[v]$ for extension by a new head value v .

Definition 2.5 (Source denotation). Denotation is structural. Variables use ρ , constants use M_s , and the Boolean forms have their ordinary meanings. From this point onward,

$$\llbracket e \rrbracket_{M_s, \rho}$$

denotes the value of e in M_s under ρ .³ The higher-order clauses are

$$\begin{aligned} \llbracket \lambda x:\tau. e \rrbracket_{M_s, \rho} &= \lambda v. \llbracket e \rrbracket_{M_s, \rho[v]}, \\ \llbracket f a \rrbracket_{M_s, \rho} &= \llbracket f \rrbracket_{M_s, \rho}(\llbracket a \rrbracket_{M_s, \rho}), \\ \llbracket \forall x:\tau. \varphi \rrbracket_{M_s, \rho} &= \forall v \in \llbracket \tau \rrbracket_B, \llbracket \varphi \rrbracket_{M_s, \rho[v]}, \\ \llbracket \exists x:\tau. \varphi \rrbracket_{M_s, \rho} &= \exists v \in \llbracket \tau \rrbracket_B, \llbracket \varphi \rrbracket_{M_s, \rho[v]}. \end{aligned}$$

Definition 2.6 (Sentences, theories, and satisfaction). A source sentence is a Boolean term in the empty context. A finite source theory is a list $\Phi = [\varphi_1, \dots, \varphi_n]$ of sentences sharing one signature Σ . We introduce the satisfaction notation

$$\begin{aligned} M_s \models \varphi &\iff \llbracket \varphi \rrbracket_{M_s, \emptyset}, \\ M_s \models^T \Phi &\iff \forall \varphi \in \Phi. M_s \models \varphi, \end{aligned}$$

where $\emptyset$ is the unique valuation of the empty context.⁴

Definition 2.7 (Source unsatisfiability). Sentence and theory unsatisfiability are semantic properties:

$$\begin{aligned} \text{Unsat}_{\text{HO}}(\varphi) &\iff \forall M_s. \neg(M_s \models \varphi), \\ \text{Unsat}_{\text{HO}}(\Phi) &\iff \forall M_s. \neg(M_s \models^T \Phi). \end{aligned}$$

For a proof-relevant model contract $\mathcal{L}(M_s)$, the restricted form is

$$\text{Unsat}_{\text{HO}}^{\mathcal{L}}(\Phi) \iff \forall M_s. \forall \ell : \mathcal{L}(M_s). \neg(M_s \models^T \Phi).$$

The Lean definitions are *Unsatisfiable*, *TheoryUnsatisfiable*, and *UnsatisfiableUnder*.⁵

3 Datatype-aware source semantics

Contents of this section. Intrinsic mutual blocks; their embedding as HO base sorts and constants; finite free values; lawful model environments.

Native datatype lowering strengthens the class of source models without changing the HO term language. Constructors, selectors, and testers remain ordinary constants in Σ ; a separate typed description says which constants they are and what laws their interpretations obey.

³Lean formalization: `Crush/Metattheory/HO/Semantics.lean`.

⁴Lean formalization: `Crush/Metattheory/HO/Semantics.lean`.

⁵Lean formalization: `Crush/Metattheory/HO/Semantics.lean`.

3.1 Intrinsic datatype blocks

Definition 3.1 (Mutual datatype block). *Fix an arity $n > 0$. A field sort is either an external base sort or a direct reference to one of the n datatypes:*

$$q ::= \text{base}(\beta) \mid \text{data}(d), \quad d \in \text{Fin}(n).$$

A field declaration pairs a descriptive name with q ; a constructor is a name and an ordered field list; a datatype declaration is a base-sort identity and an ordered constructor list. A block B assigns one datatype declaration to every $d \in \text{Fin}(n)$. Typed positions, rather than names, identify members. We write $\beta_{B,d}$ for the base-sort identity stored in the declaration at position d . After this definition we also write

$$B \vdash_D d, \quad B \vdash_C [d]K, \quad K \vdash_F F$$

for a datatype position, a constructor position in d , and a field position in K , respectively. These correspond to Lean’s scoped notations $B \vdash^D$, $B \vdash^C [d] K$, and $K \vdash^F F$.⁶

Definition 3.2 (Block well-formedness). *A block is structurally well formed when it is nonempty, its datatype base-sort identities are pairwise distinct, and every field classified as external is fresh from those datatype identities. Direct recursive fields can refer only to the current block, so unsafe cross-block cycles are not representable. The executable predicate `wellFormed` is proved equivalent to this proposition.⁷*

3.2 Finite free values

Definition 3.3 (Canonical datatype values). *For external carriers $B_0(\beta)$, the family $\text{Val}_B(B_0, d)$ and its heterogeneous constructor arguments are defined mutually. A value has exactly one finite constructor tree:*

$$v ::= K(a_1, \dots, a_m),$$

where an external field contains an element of $B_0(\beta)$ and a recursive field contains another value $v' \in \text{Val}_B(B_0, d')$. The height of a value is one plus the maximum height of its recursive fields. A block is productive when every d has a finite value after all external carriers are replaced by the unit type.⁸

Theorem 3.4 (Free-algebra laws). *Canonical values satisfy constructor injectivity and disjointness, constructor exhaustiveness, matching tester equations, matching selector equations, and strict height decrease along every recursive field. Selectors are total; on a nonmatching constructor their fallback is intentionally unconstrained.⁹*

3.3 Lawful source models

Definition 3.5 (Datatype symbol ownership). *For a block B and source signature Σ , an ownership map $\text{Symbols}(\Sigma, B)$ selects the ordinary HO constants used as each constructor, selector, and tester. The types of those constants are computed from the block: constructor fields form a curried arrow telescope, selectors map a datatype carrier to one field carrier, and testers return `Bool`.¹⁰*

⁶Lean formalization: `Crush/Metattheory/Datatype/Core.lean`.

⁷Lean formalization: `Crush/Metattheory/Datatype/Core.lean`.

⁸Lean formalization: `Crush/Metattheory/Datatype/Semantics.lean`.

⁹Lean formalization: `Crush/Metattheory/Datatype/Semantics.lean`.

¹⁰Lean formalization: `Crush/Metattheory/Datatype/Syntax.lean` and `Datatype/Model.lean`.

Definition 3.6 (Custom datatypes in the HO language). *The HO grammar has no separate datatype type former. The member d of a custom block B is represented by the ordinary source type*

$$\text{Base}(\beta_{B,d}).$$

For a field sort q , define its HO type by

$$\text{ty}_B(\text{base}(\beta)) = \text{Base}(\beta), \quad \text{ty}_B(\text{data}(d')) = \text{Base}(\beta_{B,d'}).$$

If constructor K of member d has fields $q_1, \dots, q_m$, its ownership witness selects an ordinary constant of type

$$\text{ty}_B(q_1) \rightarrow \dots \rightarrow \text{ty}_B(q_m) \rightarrow \text{Base}(\beta_{B,d}).$$

*A selector has type $\text{Base}(\beta_{B,d}) \rightarrow \text{ty}_B(q_i)$, and a tester has type $\text{Base}(\beta_{B,d}) \rightarrow \text{Bool}$. Consequently a datatype constructor, selector, or tester is not a new HO term form: it is an intrinsically typed constant application. An unowned $\text{Base}(\beta)$ remains completely opaque; ownership by a datatype environment and the lawfulness condition below are what give a selected base sort free-datatype meaning.*¹¹

Definition 3.7 (Lawful datatype model). *Let $S : \text{Symbols}(\Sigma, B)$. A witness $\text{Lawful}_B(S, M_s)$ contains an isomorphism*

$$M_s.B(\beta_{B,d}) \cong \text{Val}_B(M_s.B, d)$$

*for every datatype position d , identifies every owned constructor with its canonical curried interpretation, requires matching selectors to recover the selected field, and requires each tester to recognize exactly its constructor. Such a witness implies productivity of B .*¹²

Definition 3.8 (Datatype environment). *An environment $\mathcal{D} = [(B_1, S_1), \dots, (B_k, S_k)]$ is an ordered list of blocks and their ownership maps in one shared signature. The judgment $\text{Lawful}_{\mathcal{D}}(M_s)$ supplies a lawful witness for every entry. We write*

$$\text{Unsat}_{\text{HO}}^{\mathcal{D}}(\Phi) \iff \forall M_s. \forall \ell : \text{Lawful}_{\mathcal{D}}(M_s). \neg(M_s \models^T \Phi).$$

*The empty environment recovers ordinary source unsatisfiability.*¹³

4 The intrinsically typed first-order target

Contents of this section. Erased sorts, abstract symbol families, target terms, and target models.

Definition 4.1 (First-order sorts and erasure). *Target sorts are*

$$\kappa ::= \text{Bool} \mid \text{Base}(\beta) \mid \text{Fn}(\tau, \sigma).$$

Erasure is defined by

$$[\text{Bool}] = \text{Bool}, \quad [\text{Base}(\beta)] = \text{Base}(\beta), \quad [\tau \rightarrow \sigma] = \text{Fn}(\tau, \sigma).$$

*The last sort is opaque: it is not a function space in an arbitrary target model. Context erasure maps every entry pointwise.*¹⁴

¹¹Lean formalization: `Crush/Metatheory/Datatype/Syntax.lean`, `Datatype/Model.lean`, and `Datatype/Flattened.lean`.

¹²Lean formalization: `Crush/Metatheory/Datatype/Model.lean`.

¹³Lean formalization: `Crush/Metatheory/Datatype/Model.lean`.

¹⁴Lean formalization: `Crush/Metatheory/FO/Core.lean`.

Definition 4.2 (Declarations and symbol families). *A declaration*

$$d = (\kappa_1, \dots, \kappa_n; \kappa)$$

records an ordered argument telescope $d.\text{args} = [\kappa_1, \dots, \kappa_n]$ and a result sort $d.\text{result} = \kappa$. Let Decl be the type of all such declarations. A symbol family is a declaration-indexed family of types

$$S : \text{Decl} \longrightarrow \mathbf{Type}, \quad d \longmapsto S(d).$$

An inhabitant $p : S(d)$ is one symbol identity whose declaration is exactly d . The fiber $S(d)$ may contain no symbols, one symbol, or several distinct symbols with the same argument and result sorts. Thus the declaration is part of a symbol's type, while its identity remains separate from its declaration.¹⁵

Proposition 4.3 (FO image of datatype-owned source symbols). *Datatype membership does not change erasure: for every member d of B ,*

$$\lfloor \text{Base}(\beta_{B,d}) \rfloor = \text{Base}(\beta_{B,d}).$$

The curried HO type of an owned constructor is flattened to one FO declaration

$$(\lfloor \text{ty}_B(q_1) \rfloor, \dots, \lfloor \text{ty}_B(q_m) \rfloor; \text{Base}(\beta_{B,d})).$$

Owned selectors and testers similarly receive declarations

$$\begin{aligned} d_{\text{sel},i} &= (\text{Base}(\beta_{B,d}); \lfloor \text{ty}_B(q_i) \rfloor), \\ d_{\text{test}} &= (\text{Base}(\beta_{B,d}); \text{Bool}). \end{aligned}$$

The flattened symbol family retains these as the same source-constant identities selected by datatype ownership; only their types have been exposed as complete first-order argument telescopes. Thus a custom datatype is still a base carrier in FO, while its operations are ordinary, fully applied family symbols.¹⁶

Definition 4.4 (Typed family arguments and symbol application). *Terms and heterogeneous argument vectors are defined mutually. For a context Γ , the argument-vector family follows the declaration telescope:*

$$\begin{aligned} \text{Args}_S(\Gamma, []) &= \{()\}, \\ \text{Args}_S(\Gamma, \kappa_1 :: \bar{\kappa}) &= \{t \mid S; \Gamma \vdash t : \kappa_1\} \times \text{Args}_S(\Gamma, \bar{\kappa}). \end{aligned}$$

The only general symbol-application rule is

$$\frac{p : S(d) \quad \bar{t} : \text{Args}_S(\Gamma, d.\text{args})}{S; \Gamma \vdash p(\bar{t}) : d.\text{result}}.$$

Consequently the result sort and the complete ordered list of argument sorts are determined by the index d . A missing, extra, or incorrectly sorted argument cannot be represented. In particular, symbols are always fully applied; partial application is not a target-language operation.

For example, let $d = (\kappa_1, \kappa_2; \kappa)$ and let $p, q : S(d)$ be distinct symbols. Given $S; \Gamma \vdash u : \kappa_1$ and $S; \Gamma \vdash v : \kappa_2$, both $p(u, v)$ and $q(u, v)$ have sort κ , but they remain different terms. Neither $p(u)$ nor $p(v, u)$ is constructible.

¹⁵Lean formalization: `Crush/Metatheory/FO/Family.lean`.

¹⁶Lean formalization: `Crush/Metatheory/Datatype/Flattened.lean`.

Definition 4.5 (Target terms). *For a symbol family S , the intrinsic judgment $S; \Gamma \vdash t : \kappa$ contains variables, the fully applied family symbols above, Boolean literals and connectives, same-sort equality, and first-order quantifiers. It contains neither lambda nor a distinguished higher-order application constructor. A function value is merely an element of an opaque $\mathbf{Fn}(\tau, \sigma)$ carrier; applying that value later requires an ordinary family symbol with the appropriate declaration.*¹⁷

Why translation uses a family. The finite target syntax also exists: a finite signature is a list $\Xi = [d_0, \dots, d_m]$, and a symbol occurrence carries a typed position showing that some d_j is its declaration. That form is useful after allocation, but inconvenient during recursive translation. A recursive branch may discover a new closure or application symbol; combining two branches then concatenates their declaration traces. If terms already stored numeric signature positions, those positions would have to be weakened and reindexed after every such concatenation.

With a family, the translator instead constructs a structural identity $p : S(d)$ immediately. Recursive results concatenate an occurrence trace of existential packages (d, p) , while already constructed terms remain unchanged. For flattened defunctionalization, S is an indexed inductive family whose constructors represent source constants, application symbols, and closures. Each constructor’s result index is its exact declaration; the flattened symbol definition below spells out those shapes.

From abstract identities to concrete symbols. The abstraction is removed only at a boundary where allocation information is available. The library provides a total structural conversion to a finite signature Ξ . Write $\mathbf{Symbol}(\Xi, d)$ for the type of positions in Ξ that select declaration d . A typed resolver has one component

$$r_d : S(d) \longrightarrow \mathbf{Symbol}(\Xi, d) \quad \text{for every } d \in \text{Decl.}$$

The proved raw-SMT route can also encode family terms directly. Its encoding record similarly supplies one naming component

$$\text{name}_d : S(d) \longrightarrow \text{String} \quad \text{for every } d \in \text{Decl,}$$

with injectivity both between declarations and between two inhabitants of the same fiber, plus freshness from logical built-ins. For each traced pair (d, p) , the command encoder emits the argument and result sorts from d and the concrete identifier $\text{name}(p)$. Thus delaying allocation loses neither typing nor symbol identity.¹⁸

Definition 4.6 (Target model). *A target model M_t supplies a nonempty carrier $\llbracket \kappa \rrbracket_{M_t}$ for each target sort. A declaration denotes the curried semantic type*

$$\llbracket d \rrbracket_{M_t} = \llbracket \kappa_1 \rrbracket_{M_t} \rightarrow \dots \rightarrow \llbracket \kappa_n \rrbracket_{M_t} \rightarrow \llbracket \kappa \rrbracket_{M_t}$$

for $d = (\kappa_1, \dots, \kappa_n; \kappa)$. The model then supplies an interpretation component for every declaration:

$$I_{M_t}^d : S(d) \longrightarrow \llbracket d \rrbracket_{M_t} \quad \text{for every } d \in \text{Decl.}$$

*Hence two inhabitants of the same $S(d)$ may receive different interpretations, while every interpretation necessarily has the type prescribed by d . Target term denotation feeds the denotations of the typed argument vector to $I_{M_t}^d(p)$. The satisfaction relations $M_t \models \varphi$ and $M_t \models^T T$ are then the standard many-sorted first-order semantics.*¹⁹

¹⁷Lean formalization: `Crush/Metattheory/F0/Family.lean` and `F0/FamilySemantics.lean`.

¹⁸Lean formalization: `Crush/Metattheory/F0/Family.lean`, `Defunctionalization/TranslationResult.lean`, and `SMT/Representation.lean`.

¹⁹Lean formalization: `Crush/Metattheory/F0/FamilySemantics.lean`.

Definition 4.7 (Target-theory unsatisfiability). *A typed first-order theory is semantically unsatisfiable when it has no target model:*

$$\text{Unsat}_{\text{FO}}(T) \iff \forall M_t. \neg(M_t \models^T T).$$

5 Total flattened defunctionalization

Contents of this section. Arrow telescopes, closures, generated symbols, translation results, and the total algorithm.

5.1 Flattened arrow shape

Definition 5.1 (Arrow telescope). *Every source type has a list of leading argument types $\text{args}(\tau)$ and a non-arrow residual type $\text{res}(\tau)$:*

$$\begin{aligned} \text{args}(\text{Bool}) &= [], & \text{res}(\text{Bool}) &= \text{Bool}, \\ \text{args}(\text{Base}(\beta)) &= [], & \text{res}(\text{Base}(\beta)) &= \text{Base}(\beta), \\ \text{args}(\tau \rightarrow \sigma) &= \tau :: \text{args}(\sigma), & \text{res}(\tau \rightarrow \sigma) &= \text{res}(\sigma). \end{aligned}$$

For example, if $\alpha = \tau_1 \rightarrow \tau_2 \rightarrow v$ and v is non-arrow, then $\text{args}(\alpha) = [\tau_1, \tau_2]$ and $\text{res}(\alpha) = v$.

The indices of a target argument spine record that applying the listed arguments changes a source type α into a residual type γ . A *ground-result witness* proves that γ is Boolean or a base type; only then can the spine be emitted as one complete first-order symbol application. This indexed representation is what makes under-application impossible in the total translator. ²⁰

5.2 Captures, closures, and symbols

Definition 5.2 (Exact captures). $\text{FV}(e)$ is the duplicate-free, left-to-right list of free de Bruijn positions in e . Beneath a binder, position zero is removed and positive positions are shifted down. A closure descriptor

$$\chi = (\Gamma, \tau, \sigma, e), \quad \Sigma; \Gamma, \tau \vdash e : \sigma,$$

stores the lambda's context, domain, codomain, body, and the typed references selected from $\text{FV}(e)$. If their types are $\gamma_1, \dots, \gamma_m$, these are the exact constructor captures. ²¹

Definition 5.3 (Flattened symbol family). *The target symbol family has four semantic roles:*

1. a source constant $c : \alpha$, declared with arguments $[\text{args}(\alpha)]$ and result $[\text{res}(\alpha)]$;
2. an application symbol for every arrow $\tau \rightarrow \sigma$, declared as

$$\text{app}_{\tau, \sigma} : \text{Fn}(\tau, \sigma) \times [\text{args}(\tau \rightarrow \sigma)] \longrightarrow [\text{res}(\tau \rightarrow \sigma)];$$

3. a closure constructor

$$\text{clos}_\chi : [\gamma_1] \times \dots \times [\gamma_m] \longrightarrow \text{Fn}(\tau, \sigma);$$

The roles remain distinct even when two declarations have the same shape. Certified production hooks use the separate semantic contracts in *Hooks.lean*; the total structural translator does not add a second symbol identity for them. ²²

²⁰Lean formalization: `Crush/Metattheory/Defunctionalization/Flattened/Spine.lean`.

²¹Lean formalization: `Crush/Metattheory/Defunctionalization/Collect.lean` and `Defunctionalization/Annotate.lean`.

²²Lean formalization: `Crush/Metattheory/Defunctionalization/Flattened/Symbol.lean`.

5.3 What translation returns

Definition 5.4 (Translation result). For $\Sigma; \Gamma \vdash e : \tau$, write $\mathcal{F}[e]$ for the total flattened translation result. It contains

$$\mathcal{F}[e] = (t_e, D_e, E_e, X_e),$$

where:

- t_e is a target term in context $[\Gamma]^*$ and sort $[\tau]$;
- D_e is the ordered trace of required symbol declarations;
- E_e is the list of closure equations;
- X_e is the list of extensionality formulas;

The generated auxiliary theory, the complete target theory of one sentence, and the combined target theory of a finite source theory are

$$\begin{aligned} \text{Gen}(e) &= E_e ++ X_e, \\ \text{Theory}(\varphi) &= \text{Gen}(\varphi) ++ [t_\varphi], \\ \text{Theory}^*(\Phi) &= \text{concat}_{\varphi \in \Phi} \text{Theory}(\varphi). \end{aligned}$$

Lists are kept separate in Lean to retain provenance, then concatenated in this fixed order for semantic statements. The aggregate operation is Lean’s *translatedTheories*; every sentence is translated against the same intrinsic signature. ²³

5.4 How terms are translated

Boolean constructors, ground variables, ground constants, equality at ground types, and quantifiers translate structurally. Three cases require additional machinery.

Complete applications. The translator collects the left-associated source application spine. When its result is ground, the typed argument indices prove that all leading arrow arguments are present, so it emits exactly one flattened source or application symbol.

Residual function values. A variable, constant, or partial application of arrow type cannot be emitted as an under-applied first-order symbol. It is eta-expanded into a closure. The `LambdaBody` structure opens the complete leading lambda telescope and the theorem `denote_etaBody` proves that the resulting source function has exactly the original denotation. ²⁴

Function equality. Equality between opaque function-sort values would be stronger than pointwise equality unless extensionality is asserted. For each function equality used by the translation, X_e receives a closed formula saying that pointwise equality over the complete arrow telescope implies equality of the function values.

²³Lean formalization: `Crush/Metattheory/Defunctionalization/TranslationResult.lean` and `Defunctionalization/Flattened/Theory.lean`.

²⁴Lean formalization: `Crush/Metattheory/Defunctionalization/Flattened/Lambda.lean` and `Defunctionalization/EtaCorrectness.lean`.

Definition 5.5 (Flattened closure equation). *Suppose χ has arrow type $\alpha = \tau_1 \rightarrow \dots \rightarrow \tau_n \rightarrow v$ with ground v , exact captures $\bar{y}$, and eta-opened body b . Its defining equation has the shape*

$$\forall \bar{y} x_1 \dots x_n. \quad \text{app}_\alpha(\text{clos}_\chi(\bar{y}), x_1, \dots, x_n) = t_b,$$

where app_α abbreviates the application symbol selected by the full arrow type α . The left side uses one flattened application, not a chain of unary applications. The right side carries all obligations recursively generated while translating b .

The three mutually recursive modes—ordinary term translation, application spine translation, and lambda-body translation—use a lexicographic measure built from constructor count and a mode tag. Lean therefore accepts them as total definitions, with no partial escape hatch. ²⁵

6 The canonical model and flattened correctness

Contents of this section. Why the model is chosen; term preservation; generated validity; unsatisfiability reflection.

6.1 Modeling function values

An arbitrary first-order target model may interpret $\text{Fn}(\tau, \sigma)$ by any nonempty carrier. To prove soundness, however, it suffices to construct one target model from each source model. This is the role of the canonical model.

Definition 6.1 (Canonical flattened model). *Given a source model M_s , its canonical target model $\widehat{M}_s$ uses the same carriers at Boolean and base sorts and uses the genuine source function space at each function sort:*

$$\llbracket \text{Fn}(\tau, \sigma) \rrbracket_{\widehat{M}_s} = \llbracket \tau \rightarrow \sigma \rrbracket_{M_s}.$$

More generally, the Lean development exposes mutually inverse, type-directed maps toCanonical_τ and $\text{fromCanonical}_\tau$ so statements remain uniform at all types. In $\widehat{M}_s$:

- flattened source symbols perform repeated source application;
- $\text{app}_{\tau, \sigma}$ performs repeated source application;
- clos_χ reconstructs the exact captured source environment and denotes the source lambda.

This does not claim that all target models contain actual functions. It is a countermodel construction used only in the soundness proof. ²⁶

A source valuation ρ induces a canonical target valuation $\hat{\rho}$ by applying toCanonical pointwise.

Theorem 6.2 (Flattened term preservation). *For every source model M_s , valuation ρ , and typed source term $\Sigma; \Gamma \vdash e : \tau$.²⁷*

$$\llbracket t_e \rrbracket_{\widehat{M}_s, \hat{\rho}} = \text{toCanonical}_\tau(\llbracket e \rrbracket_{M_s, \rho}), \quad \text{where } \mathcal{F}\llbracket e \rrbracket = (t_e, D_e, E_e, X_e).$$

²⁵Lean formalization: `Crush/Metattheory/Defunctionalization/Flattened/Translate.lean`.

²⁶Lean formalization: `Crush/Metattheory/Defunctionalization/ModelExtension.lean` and `Defunctionalization/Flattened/Symbol.lean`.

²⁷Lean formalization: `Crush/Metattheory/Defunctionalization/Flattened/Denotation.lean`.

Proof architecture. One structural induction proves ordinary translation and spine translation simultaneously. Ground syntax is homomorphic. Complete spines use the currying theorem for the flattened symbol. Residual arrows use eta-closure correctness. Quantifiers use the fact that extending a canonical target valuation corresponds to extending the source valuation. $\square$

Theorem 6.3 (Validity of generated obligations). *For every e , the same canonical model satisfies the entire auxiliary theory.*²⁸

$$\widehat{M}_s \models^T \text{Gen}(e).$$

Proof architecture. A second simultaneous induction follows the three translator modes. Closure equations use term preservation and eta correctness; function extensionality is valid because the canonical function carriers are genuine source function spaces. $\square$

Theorem 6.4 (Flattened model extension). *For every source sentence φ and source model M_s .*²⁹

$$M_s \models \varphi \implies \widehat{M}_s \models^T \text{Theory}(\varphi).$$

Proof. Generated-obligation validity proves the prefix $\text{Gen}(\varphi)$. Flattened term preservation at the empty valuation proves the final assertion t_φ . $\square$

Theorem 6.5 (Aggregate flattened model extension). *For every finite source theory Φ and source model M_s ,*

$$M_s \models^T \Phi \implies \widehat{M}_s \models^T \text{Theory}^*(\Phi).$$

*All source sentences use one canonical target model and one shared flattened symbol family; the theorem is not a pointwise collection of unrelated target models.*³⁰

Proof. Induction over Φ applies the singleton model-extension theorem at the head and the induction hypothesis to the tail. Target-theory satisfaction distributes over concatenation. $\square$

Corollary 6.6 (Flattened unsatisfiability reflection). *The singleton model-extension result yields.*³¹

$$\text{Unsat}_{\text{FO}}(\text{Theory}(\varphi)) \implies \text{Unsat}_{\text{HO}}(\varphi).$$

Proof. If φ had a source model, model extension would construct a target model of $\text{Theory}(\varphi)$, contradicting the premise. $\square$

Corollary 6.7 (Aggregate flattened unsatisfiability reflection).

$$\text{Unsat}_{\text{FO}}(\text{Theory}^*(\Phi)) \implies \text{Unsat}_{\text{HO}}(\Phi).$$

*This is the theorem `target_theories_unsat_implies_source_unsat`.*³²

²⁸Lean formalization: `Crush/Metattheory/Defunctionalization/Flattened/Theory.lean`.

²⁹Lean formalization: `Crush/Metattheory/Defunctionalization/Flattened/Theory.lean`.

³⁰Lean formalization: `Crush/Metattheory/Defunctionalization/Flattened/Theory.lean`.

³¹Lean formalization: `Crush/Metattheory/Defunctionalization/Flattened/Theory.lean`.

³²Lean formalization: `Crush/Metattheory/Defunctionalization/Flattened/Theory.lean`.

6.2 The unary reference proof

The development retains a smaller classic translation. We write $\mathcal{D}[[e]]$ for this unary reference translation, which uses a binary application symbol at each arrow. A source/target logical relation is defined at Boolean, base, and arrow types; its arrow clause says that related function values produce related results when observed through unary application. The fundamental lemma proves

$$[[e]]_{M_s, \rho_s} \approx_\tau [[\mathcal{D}[[e]]]]_{M_t, \rho_t}$$

under related models and valuations. The flattened currying development then proves that the production-shaped translation and the unary reference have the same canonical denotation. Thus the unary proof is an independent semantic reference, not the endpoint of the current VCG theorem.³³

7 Raw SMT semantics

Contents of this section. Relational term evaluation; admitted commands; native datatype laws; semantic command-set equality.

The production SMT syntax is intentionally not intrinsically typed. The proof does not claim that every raw term is well sorted; instead it gives an inductive evaluation relation to exactly the fragment used by the modeled lowering pipeline.

Definition 7.1 (Raw SMT model and evaluation). *A raw model M_q contains one universe of values, a membership predicate $\text{inSort}(s, v)$, a nonemptiness witness for every concrete sort, distinguished Boolean values, literal interpretations, and a graph $\text{apply}(p, \bar{v}, w)$ for identifiers. The structure itself does not assume that this graph is functional. Global single-valuedness is the separate property **ApplyUnique**, established for the induced models used by the soundness proof; each admitted symbol declaration independently requires a total, single-valued, sort-correct graph at its declared type. The relation*

$$\text{Eval}(M_q, \eta, q, v)$$

evaluates raw term q to v under nearest-binder-first environment η . It covers bound variables, literals, logical connectives, equality, quantifiers, annotations, and user-symbol applications. A closed raw formula q holds when it evaluates to semantic truth; only now do we write $M_q \models_{\text{SMT}} q$.³⁴

Definition 7.2 (Supported commands). *The admitted command fragment contains sort and function declarations, assertions, administrative/query commands, monomorphic **declare-datatypes**, and simultaneous **define-funs-rec** blocks. Nonrecursive function and sort definitions remain outside this semantic fragment. A function declaration requires a total, single-valued, sort-correct graph. An assertion requires its formula to hold.³⁵*

Definition 7.3 (Raw native datatype semantics). *For one monomorphic **declare-datatypes** command, the predicate $\text{DatatypesHold}(M_q, B_q)$ requires:*

1. *total, typed, injective constructors and matching selector equations;*

³³Lean formalization: `Crush/Metatheory/Defunctionalization/Fundamental.lean`, `Defunctionalization/ModelExtension.lean`, and `Defunctionalization/Flattened/Currying.lean`.

³⁴Lean formalization: `Crush/Metatheory/SMT/Semantics.lean`.

³⁵Lean formalization: `Crush/Metatheory/SMT/Semantics.lean`.

2. total testers that accept their constructor and reject distinct ones;
3. disjoint constructor images and constructor exhaustiveness; and
4. a rank $r : M_q.\text{Value} \rightarrow \mathbb{N}$ that strictly decreases on every field whose sort belongs to the same mutual block.

The rank excludes cyclic or infinite raw values. Selectors are unconstrained on nonmatching constructors, exactly as in SMT-LIB.³⁶

Definition 7.4 (Raw simultaneous recursive definitions). *A supported **define-funs-rec** block has a nonempty, duplicate-free list of fresh ordinary identifiers. Each function graph is total at its declared type, and for typed arguments $\bar{v}$,*

$$\text{apply}(f, \bar{v}, w) \iff \text{Eval}(M_q, \text{reverse}(\bar{v}), q_f, w),$$

where q_f is the retained body. Every body is evaluated against the same global graph, giving simultaneous mutual-recursion semantics rather than a vacuous support marker.³⁷

Definition 7.5 (Command satisfaction and command-set equality). *The judgment $M_q \models_{\text{SMT}} C$ means that one raw model satisfies every command occurring in array C . It is deliberately insensitive to order and repeated occurrences. Define*

$$C_1 \equiv_{\text{cmd}} C_2 \iff (\forall c \in C_1. c \in C_2) \wedge (\forall c \in C_2. c \in C_1).$$

Then $C_1 \equiv_{\text{cmd}} C_2$ implies $M_q \models_{\text{SMT}} C_1 \leftrightarrow M_q \models_{\text{SMT}} C_2$. Concrete declaration-before-use ordering is checked separately by the production script validator.³⁸

Definition 7.6 (Semantic command unsatisfiability).

$$\text{Unsats}_{\text{SMT}}(C) \iff \forall M_q. \neg(M_q \models_{\text{SMT}} C).$$

This is a mathematical premise about commands, not a claim about an external solver process or its reported verdict.³⁹

8 Exact FO-to-SMT representation

Contents of this section. One shared encoding; semantic command-set representation; native datatype certificates; model lifting.

Definition 8.1 (Encoding choice). *For a typed symbol family S , an encoding $\mathcal{E}$ chooses:*

- an injective concrete SMT sort $\mathcal{E}_s(\kappa)$ for every FO sort, with **Bool** mapped to SMT **Bool**;
- an identifier $\mathcal{E}_i(p)$ for every typed symbol, injective both across declaration indices and within each symbol-family fiber, and fresh from logical built-ins;
- predicates identifying sorts and symbols owned by native commands;

³⁶Lean formalization: `Crush/Metattheory/SMT/Semantics.lean`.

³⁷Lean formalization: `Crush/Metattheory/SMT/Semantics.lean`.

³⁸Lean formalization: `Crush/Metattheory/SMT/Semantics.lean`.

³⁹Lean formalization: `Crush/Metattheory/SMT/Semantics.lean`.

- the exact native command prefix C_{native} ; and
- ordinary printable names for every nonnative symbol.

All components share this one namespace and one raw model. Native datatypes do not introduce a parallel term encoder. Given $\mathcal{E}$, we now write

$$\mathcal{A}[[t]][\mathcal{E}]$$

for the raw encoding of typed FO term t . Variables become numeric de Bruijn indices; symbols use $\mathcal{E}_i$; and logical syntax is mapped homomorphically. ⁴⁰

Definition 8.2 (Pure theory encoder). For an ordered declaration occurrence trace D and typed theory T , define

$$\text{theory}(\mathcal{E}, D, T) = C_{\text{native}} +! + C_{\text{sort}} +! + C_{\text{decl}} +! + C_{\text{assert}}.$$

Native sorts and symbols are omitted from the ordinary declaration segments. Used ordinary sorts, declarations, and assertions retain stable source order. ⁴¹

Definition 8.3 (Exact semantic theory representation). Only after defining the pure encoder and command-set equality do we introduce

$$\text{Rep}(\mathcal{E}, T, C) \iff \exists D. C \equiv_{\text{cmd}} \text{theory}(\mathcal{E}, D, T).$$

This is *TheoryRepresentation*. It is intentionally weaker than array equality and exactly as strong as the raw command semantics: production may interleave independent declarations and assertions or suppress duplicates, but it may neither omit nor add a semantic obligation. ⁴²

Theorem 8.4 (Exact aggregate representation). Let D_Φ concatenate every declaration occurrence produced while translating the finite source theory Φ . The aggregate encoder is

$$\text{Commands}(\mathcal{E}, \Phi) = \text{theory}(\mathcal{E}, D_\Phi, \text{Theory}^*(\Phi)),$$

and

$$\text{Rep}(\mathcal{E}, \text{Theory}^*(\Phi), \text{Commands}(\mathcal{E}, \Phi)).$$

The theorem is *encode_theories*; *encode_translation* is its singleton counterpart. ⁴³

8.1 Native datatype representation

Definition 8.5 (One represented native block). For a source block B , ownership map S , shared encoding $\mathcal{E}$, and allocated raw block encoding A , write $A_s(d)$ for the concrete SMT sort name allocated to member d . A $\text{Representation}(B, S, \mathcal{E}, A)$ proves:

- the exact emitted *declare-datatypes* command is well formed;
- constructor, selector, and tester ownership is exclusive;
- every datatype sort and symbol is marked native; and

⁴⁰Lean formalization: `Crush/Metattheory/SMT/Representation.lean`.

⁴¹Lean formalization: `Crush/Metattheory/SMT/Representation.lean`.

⁴²Lean formalization: `Crush/Metattheory/SMT/Representation.lean`.

⁴³Lean formalization: `Crush/Metattheory/SMT/Representation.lean`.

- all encoded sorts and identifiers equal the names in A , including indexed tester identifiers.

An environment representation lists such certificates in dependency order and proves that their command array is exactly C_{native} .⁴⁴

Definition 8.6 (One custom datatype across HO, FO, and SMT). For an accepted ground Lean datatype instance, structural reification chooses a block member d and its identity $\beta_{B,d}$. The sort component of the lowering pipeline is

$$\begin{array}{ccc} \boxed{\text{ground Lean datatype instance}} & \xrightarrow{\text{structural reification}} & \boxed{\text{Base}(\beta_{B,d}) \text{ in HO}} \\ \boxed{\text{Base}(\beta_{B,d}) \text{ in HO}} & \xrightarrow{\text{erase}} & \boxed{\text{Base}(\beta_{B,d}) \text{ in FO}} \xrightarrow{\mathcal{E}_s} \boxed{A_s(d) \text{ in raw SMT}} \end{array}$$

The first arrow records declaration structure and ownership; it is not a denotational semantics for arbitrary *Lean.Expr*. The semantic bridge begins with $\text{Lawful}_B(S, M_s)$, which identifies the HO base carrier with the canonical finite value family $\text{Val}_B(M_s.B, d)$.

The operation component follows the same ownership map. An owned HO constructor, selector, or tester becomes its declaration-indexed FO source symbol, and $\mathcal{E}_i$ maps that symbol to the corresponding constructor, selector, or indexed tester identifier recorded by A . The native **declare-datatypes** command declares those identifiers together, so the same object is not also emitted as an ordinary uninterpreted sort or function. All other HO and FO syntax remains generic: datatype-specific meaning enters only through ownership, lawfulness, and the native representation certificate.⁴⁵

Theorem 8.7 (Native command validity). If M_s is lawful for the represented datatype environment $\mathcal{D}$, then the raw model induced by $\mathcal{E}$ and the canonical FO model satisfies the complete native command prefix:

$$\text{Lawful}_{\mathcal{D}}(M_s) \implies M_q(\mathcal{E}, \widehat{M}_s) \models_{\text{SMT}} C_{\text{native}}.$$

The proof establishes constructor typing and injectivity, selector/tester laws, exhaustiveness, disjointness, and rank decrease in the shared raw universe. The dependency-extension and carry modules prove that each earlier block remains valid after every later disjoint block is installed.⁴⁶

8.2 One induced raw model

Definition 8.8 (Induced raw model). Given $\mathcal{E}$ and a typed target model M_t , the induced raw universe tags each value by its intrinsic FO sort. Sort membership checks that tag. The graph for an encoded symbol decodes typed arguments, applies M_t 's curried interpretation, and retags the result. An optional **ExtraGraph** adds certified derived identifiers only when they are disjoint from all identifiers owned by $\mathcal{E}$.⁴⁷

Theorem 8.9 (Complete representation soundness). Let $\mathcal{D}$ be the datatype environment represented by $\mathcal{E}$. If

$$\text{Rep}(\mathcal{E}, T, C), \quad \text{Lawful}_{\mathcal{D}}(M_s), \quad \widehat{M}_s \models^T T,$$

then there exists one raw model M_q with $M_q \models_{\text{SMT}} C$. Native datatype commands, ordinary sort declarations, ordinary symbol declarations, and all assertions are validated in that same model.⁴⁸

⁴⁴Lean formalization: `Crush/Metatheory/SMT/DatatypeWF.lean` and `SMT/DatatypeRepresentation.lean`.

⁴⁵Lean formalization: `Crush/Metatheory/Reification/Datatype.lean`, `Datatype/Flattened.lean`, `SMT/Datatype.lean`, and `SMT/DatatypeRepresentation.lean`.

⁴⁶Lean formalization: `Crush/Metatheory/SMT/DatatypeCanonical.lean`, `SMT/DatatypeLifted.lean`, and `SMT/DatatypeCarry.lean`.

⁴⁷Lean formalization: `Crush/Metatheory/SMT/Model.lean` and `SMT/ModelExtension.lean`.

⁴⁸Lean formalization: `Crush/Metatheory/SMT/Soundness.lean`.

Model construction. The term-evaluation theorem relates $\mathcal{A}[[t]][\mathcal{E}]$ to intrinsic denotation in M_t . Separate lemmas validate ordinary declarations and assertions. Native validity supplies the exact prefix. Finally, $C \equiv_{\text{cmd}} \text{theory}(\mathcal{E}, D, T)$ transports satisfaction to the represented production-shaped array. $\square$

Corollary 8.10 (Aggregate lowering reflection). *If*

$$\text{Rep}(\mathcal{E}, \text{Theory}^*(\Phi), C) \quad \text{and} \quad \text{Unsat}_{\text{SMT}}(C),$$

then

$$\text{Unsat}_{\text{HO}}^{\mathcal{D}}(\Phi).$$

This is `commands_unsat_implies_source_theory_unsat`. It composes complete representation soundness with aggregate flattened model extension. ⁴⁹

9 Guarded carriers and guarded SMT encoding

Contents of this section. Carrier equivalences; related FO models; recursive datatype guards; one guarded raw-model theorem.

Some source carriers occupy a proper subset of the corresponding production SMT carrier. The primary external example is $\mathbb{N}$ represented in $\mathbb{Z}$ by the predicate $0 \leq z$. Datatype values recursively inherit guards from their external fields.

9.1 Guarded carrier equivalence

Definition 9.1 (Guarded relation). *A guarded relation $R : S \rightsquigarrow T$ contains*

$$\text{encode} : S \rightarrow T, \quad \text{guard} : T \rightarrow \text{Prop}, \quad \text{decode} : (t : T) \rightarrow \text{guard}(t) \rightarrow S,$$

together with nonemptiness of S and the laws

$$\text{guard}(\text{encode}(s)), \quad \text{decode}(\text{encode}(s)) = s, \quad \text{guard}(t) \Rightarrow \text{encode}(\text{decode}(t)) = t.$$

Thus S is equivalent to the guarded subtype of T . Encoding is injective; source equality is preserved and reflected; and source quantifiers correspond to the following guarded target quantifiers. ⁵⁰

$$\forall t : T. \text{guard}(t) \Rightarrow P(t), \quad \exists t : T. \text{guard}(t) \wedge P(t).$$

Definition 9.2 (FO carrier and model relation). *A carrier relation R assigns a guarded relation to every nonlogical FO carrier; Boolean propositions use the identity relation. For source and target family models M_0, M_1 , a model relation $M_0 \sim_R M_1$ requires every typed symbol interpretation to map encoded arguments to the encoded source result. This relational form admits both the generic pointwise lifting and native datatype interpretations.* ⁵¹

Definition 9.3 (Guard-restricted FO denotation). *For target model M_1 , relation R , and lifted valuation, guarded denotation uses ordinary semantics for variables, symbols, connectives, and equality, but restricts both quantifiers to the selected carrier guard. It is written explicitly as*

$$\text{guardDenote}(e, M_1, R, \nu).$$

No bracket notation is introduced because ordinary denotation and guarded denotation must remain visually distinct. ⁵²

⁴⁹Lean formalization: `Crush/Metattheory/SMT/Soundness.lean`.

⁵⁰Lean formalization: `Crush/Metattheory/Guarded/Encoding.lean`.

⁵¹Lean formalization: `Crush/Metattheory/FO/Guarded.lean`.

⁵²Lean formalization: `Crush/Metattheory/FO/Guarded.lean`.

Theorem 9.4 (FO guarded preservation). *If $M_0 \sim_R M_1$, then for every typed term e and source valuation ν ,*

$$\text{guardDenote}(e, M_1, R, \text{lift}_R(\nu)) = R_\kappa.\text{encode}(\llbracket e \rrbracket_{M_0, \nu}),$$

where κ is the result sort of e . The theorem handles variables, symbols, all logical forms, equality, and both quantifiers uniformly. ⁵³

9.2 Datatype guard lifting

Definition 9.5 (Structural datatype guard). *Let external fields be related pointwise by R_0 . A target datatype value is well formed when every external payload satisfies its base guard and every recursive payload is structurally well formed. This predicate is $\text{WF}_B(R_0, v)$. Encoding maps the constructor tree recursively and encoding or decoding is inverse exactly on WF_B values. Hence every productive block lifts R_0 to a guarded relation on each of its datatype carriers.* ⁵⁴

Theorem 9.6 (Production-shaped guard equivalence). *Production defines a predicate wf_T as a conjunction over every constructor tester and its selectors: when a tester accepts, every matching selected field must satisfy its guard. For a productive block,*

$$\text{WF}_B(R_0, v) \iff \text{SelWF}_B(R_0, v).$$

This is Val.wf_iff_selWF . The enlarged carrier model installs native constructors, selectors, and testers while preserving the same per-symbol model relation. ⁵⁵

9.3 Guard-aware raw syntax

Definition 9.7 (Guarding policy and guarded encoder). *A guarding policy $\mathcal{G} = (\mathcal{E}, g)$ reuses one ordinary encoding $\mathcal{E}$ and optionally supplies a raw guard term $g(\kappa, q)$ for a raw term q of sort κ . We now write*

$$\mathcal{G}[\llbracket e \rrbracket][\mathcal{G}]$$

for the shared term encoder instantiated with g . At a universal binder it emits guard implication; at an existential binder it emits guard conjunction. Every other term constructor is encoded by the same function used for $\mathcal{A}[\llbracket e \rrbracket][\mathcal{E}]$. ⁵⁶

Definition 9.8 (Guarded theory representation). *Let C_{derived} be an exact certified command segment, such as the mutual wf_T definitions. Define the following guarded command sequence and representation.* ⁵⁷

$$\text{gtheory}(\mathcal{G}, C_{\text{derived}}, D, T) = C_{\text{native}} +!+ C_{\text{derived}} +!+ C_{\text{ordinary}} +!+ C_{\text{gassert}}.$$

Then

$$\text{Rep}_g(\mathcal{G}, C_{\text{derived}}, T, C) \iff \exists D. C \equiv_{\text{cmd}} \text{gtheory}(\mathcal{G}, C_{\text{derived}}, D, T).$$

Definition 9.9 (Semantic guard contract). *For the induced raw model, target model M_1 , and carrier relation R , the contract $\text{GuardSem}(\mathcal{G}, M_1, R)$ says:*

⁵³Lean formalization: `Crush/Metattheory/F0/Guarded.lean`.

⁵⁴Lean formalization: `Crush/Metattheory/Datatype/Guarded.lean`.

⁵⁵Lean formalization: `Crush/Metattheory/Datatype/Guarded.lean`, `Datatype/Carrier.lean`, `Datatype/FamilyModel.lean`, and `Datatype/Flattened.lean`.

⁵⁶Lean formalization: `Crush/Metattheory/SMT/Guarded.lean`.

⁵⁷Lean formalization: `Crush/Metattheory/SMT/Guarded.lean`.

- if syntax is omitted for κ , every target value satisfies $R_\kappa.\text{guard}$; and
- if $g(\kappa, q)$ is emitted and q evaluates to v , the guard term evaluates exactly to $R_\kappa.\text{guard}(v)$.

The compositional *TermSemantics* form applies the same statement to an arbitrary already-evaluated raw term, which is needed for datatype selectors. ⁵⁸

Theorem 9.10 (Guarded raw term preservation). *If $M_0 \sim_R M_1$ and the guarding policy has the semantic guard contract, then the raw term $\mathcal{G}[[e]]\mathcal{G}$ evaluates to the tagged encoding of $[[e]]_{M_0, \nu}$. This is `guardTerm_rel_eval`; it composes FO guarded preservation with the raw evaluator.* ⁵⁹

9.4 Native blocks and recursive guards in one model

The datatype target is built by a dependency fold. At each step, the current block receives its enlarged free-algebra carrier and native interpretations. An `Ordered` witness proves every earlier native command is disjoint from every later block, not only from its immediate successor. The carry theorems transport constructor, selector, tester, exhaustive, and rank laws through the entire suffix.

Each represented block has one exact `define-funs-rec` guard command. Its retained names and binder determine the same shared `wfBody` syntax used by production. The raw evaluator proves the Boolean conjunction and field clauses, then `wfDefs_valid` proves the simultaneous command graph. The exact guard command array is indexed by the same represented block list as the native prefix. ⁶⁰

Theorem 9.11 (Complete guarded model lifting). *Assume*

$$\text{Rep}_g(\mathcal{G}, C_{\text{derived}}, T, C), \quad M_0 \sim_R M_1, \quad M_0 \models^T T,$$

the semantic guard contract, validity of C_{native} , and validity of C_{derived} . Then one raw model satisfies C . This is `guarded_lift`; native declarations, derived recursive definitions, ordinary declarations, and guarded assertions all share one raw graph. ⁶¹

Remark 9.12 (Why the uniform guard interpretation matters). A model-indexed target construction may be useful as a low-level premise, but placing it inside the definition of source unsatisfiability would silently restrict the quantified source-model class. The production-facing uniform interpretation instead supplies a target construction for every datatype-lawful source model. The resulting reflection theorem concludes $\text{Unsat}_{\text{HO}}^D$, with only intrinsic datatype lawfulness in the quantified model class. ⁶²

10 Lean reification and intrinsic VCG

Contents of this section. Partial structural reification; common-environment theories; pure and stateful aggregate lowering.

⁵⁸Lean formalization: `Crush/Metatheory/SMT/GuardedSoundness.lean`.

⁵⁹Lean formalization: `Crush/Metatheory/SMT/GuardedSoundness.lean`.

⁶⁰Lean formalization: `Crush/Metatheory/SMT/Datatype.lean`, `SMT/DatatypeGuard.lean`, `SMT/DatatypeGuarded.lean`, and `SMT/DatatypeCarry.lean`.

⁶¹Lean formalization: `Crush/Metatheory/SMT/GuardedSoundness.lean`.

⁶²Lean formalization: `Crush/Metatheory/VCG/Datatype.lean`.

10.1 The supported Lean boundary

Definition 10.1 (Structural reification witness). *The executable reifier recognizes supported normalized Lean-expression constructors and returns an existentially typed HO term. A witness retains the recognized expression shape, the intrinsic term shape, their executable correspondence proof, and the exact typed context, ordinary-signature, and datatype bridges used by reification. Lean’s inferred type is checked at every accepted node.* ⁶³

The witness is structural. It does not define a denotation for arbitrary Lean expressions and therefore does not by itself prove that the original Lean proposition and its intrinsic sentence have equal truth values. Unsupported syntax returns no witness.

Definition 10.2 (Reified finite theory). *Given the exact ordered expression list*

$$[E_1, \dots, E_n],$$

common-environment reification first selects one datatype environment $\mathcal{D}$ and one ordinary signature tail. It then retains

$$(E_1, \varphi_1), \dots, (E_n, \varphi_n)$$

pointwise in the same order, with every φ_i indexed by the same complete signature. The type `ReifiedSentencesFor` prevents omission, reordering, or substitution of a fact. The partial executable function `reifyTheory?` returns this package only when every fact reifies. ⁶⁴

Definition 10.3 (Supported datatype reification). *Datatype discovery produces dependency-ordered, monomorphic, productive ground blocks. It rejects indices, dependent or function-valued fields, non-productivity, unsafe indirect recursion, recursors, and quotient primitives. Discovery is read-only; name allocation and command emission begin only after the complete structural plan is known.* ⁶⁵

10.2 Pure aggregate VCG

Definition 10.4 (Pure VCG commands). *For encoding $\mathcal{E}$ and finite intrinsic source theory Φ , the pure VCG function computes*

$$\text{theoryCommands}(\mathcal{E}, \Phi) = \text{Commands}(\mathcal{E}, \Phi)$$

and retains

$$\text{Rep}(\mathcal{E}, \text{Theory}^*(\Phi), \text{theoryCommands}(\mathcal{E}, \Phi)).$$

For guarding policy $\mathcal{G}$ and exact derived segment C_{derived} , `guardedTheoryCommands` constructs the corresponding `Repg` witness. Sentence-level functions are singleton conveniences. ⁶⁶

Definition 10.5 (Stateful aggregate VCG). *The aggregate runner begins with fresh translation bookkeeping and installs the pure aggregate command array in the state’s command field. Its representation witness retains both exact command equality and the semantic theory representation. The guarded aggregate runner does the same for the guarded array and certified derived segment. Freshness prevents either route from inheriting commands or trust markers from a previous direct run. Lean names these runners `runTheory` and `runGuardedTheory`, and names the witness `TheoryStateRepresents`.* ⁶⁷

⁶³Lean formalization: `Crush/Metattheory/Reification/Type.lean`, `Reification/Term.lean`, `Reification/Capture.lean`, `Reification/Reify.lean`, and `Reification/Witness.lean`.

⁶⁴Lean formalization: `Crush/Metattheory/Reification/Witness.lean`.

⁶⁵Lean formalization: `Crush/Metattheory/Reification/Datatype.lean`.

⁶⁶Lean formalization: `Crush/Metattheory/VCG/Generate.lean`.

⁶⁷Lean formalization: `Crush/Metattheory/VCG/Stateful.lean`.

Theorem 10.6 (Aggregate intrinsic VCG soundness). *Let C be the command array in a fresh aggregate VCG state, let $\mathcal{D}$ be represented by the same encoding, and assume $\text{Unsat}_{\text{SMT}}(C)$. Then*

$$\text{Unsat}_{\text{HO}}^{\mathcal{D}}(\Phi).$$

This is `TheoryStateRepresents.unsat_source`; `runTheory_unsat_implies_source_unsat` is its executable specialization. The guarded finite-theory theorem uses one uniform `GuardInterpretation` and reaches the same source-model class. ⁶⁸

11 Production certificates and whole-theory agreement

Contents of this section. Exact final-state traces; static representation evidence; proof-producing command-set comparison; production reflection.

The intrinsic VCG constructs a fresh canonical array. The live production translator instead interleaves declarations and assertions, removes duplicate occurrences, attaches provenance annotations to assertions, and allows extension handlers. The production layer therefore uses semantic command-set agreement, not an unattainable array equality.

11.1 Exact final-state component traces

Definition 11.1 (Certified production datatype environment). *A `CertifiedDataEnv` retains:*

- *the exact source expression, datatype environment, ordinary signature bridge, and optional reified sentence indexed by that environment;*
- *the complete final production command snapshot C_{prod} ;*
- *a typed native-command trace; and*
- *a typed recursive-guard trace.*

Each trace entry records an index in the final snapshot and proves that the command at that index is the exact command belonging to the intrinsically indexed block. Traces are rebuilt after all facts are emitted; no intermediate prefix is presented as final evidence. ⁶⁹

Definition 11.2 (Production representation evidence). *For a certificate P :*

- *$P.\text{Representation}(\mathcal{E})$ links every typed native trace entry to one shared encoding and proves that its exact native prefix is the certified command list;*
- *$P.\text{GuardRepresentation}(\mathcal{G})$ links the recursive guard trace to fresh, injective identifiers independently of any model; and*
- *$P.\text{GuardInterpretation}(\mathcal{G})$ realizes that static guard representation for every datatype-lawful source model.*

These are premises of the production refinement boundary. They are not a second term encoder or a second datatype soundness theorem. ⁷⁰

⁶⁸Lean formalization: `Crush/Metattheory/VCG/Soundness.lean`.

⁶⁹Lean formalization: `Crush/Metattheory/VCG/Datatype.lean`.

⁷⁰Lean formalization: `Crush/Metattheory/VCG/Datatype.lean`.

11.2 Whole-theory comparison

Definition 11.3 (Semantically transparent normalization). *Production assertions may have a root `:named` annotation, while the intrinsic encoder emits the unannotated body. A proved normalizer removes only that root annotation, and its theorem proves command satisfaction unchanged. A leading `set-logic` command is likewise proved semantically neutral. No other command is erased. The Lean normalizer is `stripAssertionAnnotation`.* ⁷¹

Definition 11.4 (Whole-theory agreement). *Let W be an exact `ReifiedSentencesFor` witness with intrinsic theory Φ_W . A *TheoryAgreement* is the proposition*

$$\text{Rep}_g(\mathcal{G}, C_{\text{guard}}, \text{Theory}^*(\Phi_W), \text{strip}(C_{\text{prod}})).$$

The executable `TheoryAgreement.build?` computes the expected guarded aggregate commands and checks mutual command inclusion. It returns proof-carrying agreement only when the Boolean comparison is proved equivalent to $\equiv_{\text{cmd}}$. Thus reordering and duplicate suppression are accepted, while a missing or extra semantic obligation is rejected. ⁷²

Theorem 11.5 (Production whole-theory reflection). *Let P be a production certificate with native representation, static guard representation, uniform guard interpretation, and whole-theory agreement for W . If the exact completed snapshot is semantically unsatisfiable, then*

$$\text{Unsat}_{\text{SMT}}(C_{\text{prod}}) \implies \text{Unsat}_{\text{HO}}^{\mathcal{D}}(\Phi_W).$$

The theorem `TheoryAgreement.unsat_source_script` also admits the leading logic-selection command returned by `buildScript`. `SingleFactAgreement` is defined only as the singleton specialization of this aggregate path. ⁷³

11.3 Operational trust status

`TranslateState.status` is proved exactly when its trust-reason array is empty; otherwise it retains the nonempty reasons in a `trusted` constructor. This operational classification is not itself a semantic representation witness. The legacy extensible `emitTerm` route is marked direct-trusted at its root, and unrestricted handlers, uncertified constants or closures, and native-HO emission retain additional reasons. Certified local components become a whole-run theorem only through the exact aggregate agreement described above. ⁷⁴

12 Proof architecture and exact scope

Contents of this section. The composed aggregate theorem; modeled production behavior; explicit external boundaries.

Figure 1 summarizes the dependency chain. Solid arrows are backed by Lean theorems; the dashed reification arrow is structural, not denotational.

⁷¹Lean formalization: `Crush/Metattheory/VCG/Production.lean`.

⁷²Lean formalization: `Crush/Metattheory/VCG/Production.lean`.

⁷³Lean formalization: `Crush/Metattheory/VCG/Production.lean`.

⁷⁴Lean source: `Crush/Metattheory/VCG/Trust.lean`, `Crush/Translation/Monad.lean`, `Crush/Metattheory/VCG/Stateful.lean`, and `Crush/Metattheory/VCG/Production.lean`.




Figure 1: The current aggregate proof architecture. Datatype lawfulness and a uniform guard interpretation parameterize the model construction. Command-set agreement models production reordering and duplicate suppression exactly at the membership-based raw semantics.

12.1 What the composed theorem proves

1. HO and FO syntax is intrinsically typed, and flattened defunctionalization is total on the complete modeled core.
2. Every source fact contributes its translated body, closure equations, extensionality obligations, and declaration occurrences to one combined target theory over one signature.
3. One canonical target model simultaneously satisfies that whole target theory whenever one source model satisfies the whole source theory.
4. Productive monomorphic datatype blocks have finite free-algebra source semantics and exact native raw-command semantics.
5. Guarded carrier relations preserve typed terms, equality, and both quantifiers; recursive datatype guards denote exactly the structural image of the source carrier.
6. One induced raw model validates native datatype commands, simultaneous guard definitions, ordinary declarations, and every guarded assertion.
7. Exact representation is semantic command-set equivalence. It therefore admits production interleaving and duplicate elimination without weakening or strengthening the model obligations.
8. Common-environment reification retains the exact ordered Lean fact array, and fresh aggregate VCG retains an exact representation theorem.
9. Production traces reconnect native and recursive-guard commands to the completed final snapshot. `TheoryAgreement.build?` checks the remaining whole-array command set and returns proof-producing evidence.
10. Semantic unsatisfiability of the agreed completed snapshot reflects to unsatisfiability of the entire intrinsic source theory in every datatype-lawful source model.

12.2 Exact boundaries

- No denotational semantics is postulated for arbitrary `Lean.Expr`; the final proposition concerns the reified intrinsic theory.
- Semantic command unsatisfiability is a premise. Correctness of an external solver’s reported `unsat` answer requires certificate replay or another independently checked argument.
- The live extensible translator does not automatically construct the shared encoding, native representation, static guard representation, uniform guard interpretation, or common aggregate reification witness. Production reflection applies once those proof objects and exact agreement are present.
- `TranslateState.status` is operational metadata; it cannot replace a `TheoryRepresentation` or `TheoryAgreement` witness.
- The live allocator is finite, while the intrinsic `SMT.Encoding` interface supplies total injections over the modeled symbol family. The current theorem takes that total encoding as an explicit representation premise rather than claiming it is extracted from every allocator state.
- Unsupported dependent syntax, indexed or nonproductive inductives, dependent or function-valued datatype fields, unsafe indirect recursion, unrestricted translation handlers, nonrecursive raw definitions, and native-HO commands remain outside the modeled theorem.

13 Machine-checked theorem inventory

Contents of this section. The named results that implement the proof story.

The principal declarations, in proof order, are:

1. `Datatype.ctor_inj`, `Datatype.ctor_ne`, `Datatype.ctor_cases`, `Datatype.sel_ctor`, `Datatype.test_ctor`, and `Datatype.Args.get_height_lt_ctor`: finite free-algebra datatype laws.
2. `Datatype.Lawful.productive`: any lawful source datatype carrier supplies a productive intrinsic block.
3. `Flattened.LambdaBody.denote_etaBody` and `Flattened.TargetArguments.completeApp_denote`: eta opening and one flattened application preserve repeated source application.
4. `Flattened.translate_denote`: the translated term component has the source denotation in the canonical model.
5. `Flattened.generated_valid`: all generated closure equations and extensionality formulas are valid in that same model.
6. `Flattened.model_extension_theory`: one source model satisfying a finite theory extends to the combined flattened target theory.
7. `Flattened.target_theories_unsat_implies_source_unsat`: combined FO unsatisfiability reflects to the complete finite HO theory.
8. `SMT.encode_theories`: the pure encoder gives an exact semantic command-set representation of the combined theory.
9. `SMT.Datatype.command_sound` and `SMT.Datatype.EnvRepresentation.native_valid`: one block and the complete dependency-ordered native prefix satisfy the raw datatype laws.
10. `FO.FamilyTerm.guardDenote_rel`: related FO models preserve every term under guard-restricted quantifier semantics.

11. `Datatype.Val.wf_iff_selWF`: the structural datatype image equals the tester/selector predicate emitted by production.
12. `SMT.Datatype.wfDefs_valid`: exact mutual `wf_T` definitions satisfy the raw simultaneous graph semantics.
13. `SMT.guardTerm_rel_eval`: guard-aware raw evaluation agrees with denotation in any symbol-related target model.
14. `SMT.guarded_lift`: native commands, derived definitions, ordinary declarations, and guarded assertions hold in one raw model.
15. `SMT.representation_sound`: the unguarded complete representation has the corresponding one-model lifting theorem.
16. `SMT.commands_unsat_implies_source_theory_unsat`: semantic raw-command unsatisfiability reflects to the lawful intrinsic HO theory.
17. `Reification.reifyTheory?`: an exact Lean fact array reifies under one datatype environment and signature bridge without reordering.
18. `VCG.theoryCommands_represents` and `VCG.guardedTheoryCommands_represents`: pure aggregate generation retains ordinary and guarded representation evidence.
19. `VCG.runTheory_represents` and `VCG.runGuardedTheory_represents`: fresh stateful runs retain those exact aggregate representations.
20. `VCG.TheoryStateRepresents.unsat_source` and `VCG.runGuardedTheory_unsat_implies_source_unsat`: the intrinsic stateful compositions reflect to source-theory unsatisfiability.
21. `VCG.CertifiedDataEnv.Representation.theory_unsat`: one production datatype and guard package reflects any exactly represented finite source theory.
22. `VCG.TheoryAgreement.unsat_source` and `VCG.TheoryAgreement.unsat_source_script`: the complete production snapshot, with optional leading logic command, reflects through checked whole-theory agreement.

The earlier unary results—`fundamental`, `fundamental_sentence`, and the unary `target_unsat_implies_source_unsat`—remain machine-checked as an independent reference model.